

Parallel and Distributed Computing

-- Architecture

Architectural styles

Basic idea

A style is formulated in terms of


- ▶ (replaceable) components with well-defined interfaces
- ▶ the way that components are connected to each other
- ▶ the data exchanged between components
- ▶ how these components and connectors are jointly configured into a system.

Connector

A mechanism that mediates communication, coordination, or cooperation among components. **Example**: facilities for (remote) procedure call, messaging, or streaming.


Layered architecture

Different layered organizations



Example: communication protocols

Protocol, service, interface



Two-party communication

Server

```
1 from socket import *
2 s = socket(AF_INET, SOCK_STREAM)
3 (conn, addr) = s.accept() # returns new socket and addr. client
4 while True: # forever
5     data = conn.recv(1024) # receive data from client
6     if not data: break # stop if client stopped
7     conn.send(str(data)+"*") # return sent data plus an "*"
8     conn.close() # close the connection
```

Client

```
1 from socket import *
2 s = socket(AF_INET, SOCK_STREAM)
3 s.connect((HOST, PORT)) # connect to server (block until accepted)
4 s.send('Hello, world') # send some data
5 data = s.recv(1024) # receive the response
6 print data # print the result
7 s.close() # close the connection
```

Application Layering

Traditional three-layered view

- ▶ **Application-interface layer** contains units for interfacing to users or external applications
- ▶ **Processing layer** contains the functions of an application, i.e., without specific data
- ▶ **Data layer** contains the data that a client wants to manipulate through the application components

Application Layering

Traditional three-layered view


- ▶ **Application-interface layer** contains units for interfacing to users or external applications
- ▶ **Processing layer** contains the functions of an application, i.e., without specific data
- ▶ **Data layer** contains the data that a client wants to manipulate through the application components

Observation

This layering is found in many distributed information systems, using traditional database technology and accompanying applications.

Application Layering


Example: a simple search engine



Object-based style

Essence

Components are objects, connected to each other through procedure calls. Objects may be placed on different machines; calls can thus execute across a network.



Encapsulation

Objects are said to **encapsulate data** and offer **methods on that data** without revealing the internal implementation.

RESTful architectures

Essence

View a distributed system as a collection of resources, individually managed by components. Resources may be added, removed, retrieved, and modified by (remote) applications.

1. Resources are identified through a single naming scheme
2. All services offer the same interface
3. Messages sent to or from a service are fully self-described
4. After executing an operation at a service, that component forgets everything about the caller

Basic operations

Operation	Description
PUT	Create a new resource
GET	Retrieve the state of a resource in some representation
DELETE	Delete a resource
POST	Modify a resource by transferring a new state

Example: Amazon's Simple Storage Service

Essence

Objects (i.e., files) are placed into **buckets** (i.e., directories). Buckets cannot be placed into buckets. Operations on ObjectName in bucket BucketName require the following identifier:

<http://BucketName.s3.amazonaws.com/ObjectName>

Typical operations

All operations are carried out by sending HTTP requests:

- ▶ Create a bucket/object: PUT, along with the URI
- ▶ Listing objects: GET on a bucket name
- ▶ Reading an object: GET on a full URI

On interfaces

Issue

Many people like RESTful approaches because the interface to a service is so simple. The catch is that much needs to be done in the [parameter space](#).

Amazon S3 SOAP interface

Bucket operations	Object operations
ListAllMyBuckets	PutObjectInline
CreateBucket	PutObject
DeleteBucket	CopyObject
ListBucket	GetObject
GetBucketAccessControlPolicy	GetObjectExtended
SetBucketAccessControlPolicy	DeleteObject
GetBucketLoggingStatus	GetObjectAccessControlPolicy
SetBucketLoggingStatus	SetObjectAccessControlPolicy

On interfaces

Simplifications

Assume an interface `bucket` offering an operation `create`, requiring an input string such as `mybucket`, for creating a bucket “mybucket.”

On interfaces

Simplifications

Assume an interface `bucket` offering an operation `create`, requiring an input string such as `mybucket`, for creating a bucket “mybucket.”

SOAP

```
import bucket  
bucket.create("mybucket")
```

On interfaces

Simplifications

Assume an interface `bucket` offering an operation `create`, requiring an input string such as `mybucket`, for creating a bucket “mybucket.”

SOAP

```
import bucket  
bucket.create("mybucket")
```

RESTful

```
PUT "http://mybucket.s3.amazonaws.com/"
```

On interfaces

Simplifications

Assume an interface `bucket` offering an operation `create`, requiring an input string such as `mybucket`, for creating a bucket “mybucket.”

SOAP

```
import bucket  
bucket.create("mybucket")
```

RESTful

```
PUT "http://mybucket.s3.amazonaws.com/"
```

Conclusions



Are there any to draw?

Coordination

Temporal and referential coupling

	Temporally coupled	Temporally decoupled
Referentially coupled	Direct	Mailbox
Referentially decoupled	Event-based	Shared data space

Event-based and Shared data space



Using legacy to build middleware

Problem



The interfaces offered by a legacy component are most likely not suitable for all applications.

Solution

A **wrapper** or **adapter** offers an interface acceptable to a client application. Its functions are transformed into those available at the component.

Organizing wrappers

Two solutions: 1-on-1 or through a broker



Complexity with N applications


- ▶ **1-on-1**: requires $N \times (N - 1) = O(N^2)$ wrappers
- ▶ **broker**: requires $2N = O(N)$ wrappers

Developing adaptable middleware

Problem

Middleware contains solutions that are good for **most** applications $\Rightarrow$ you may want to adapt its behavior for specific applications.

Intercept the usual flow of control




Centralized system architectures

Basic Client–Server Model

Characteristics:

- ▶ There are processes offering services (**servers**)
- ▶ There are processes that use services (**clients**)
- ▶ Clients and servers can be on different machines
- ▶ Clients follow request/reply model with respect to using services




Multi-tiered centralized system architectures

Some traditional organizations


- ▶ **Single-tiered:** dumb terminal/mainframe configuration
- ▶ **Two-tiered:** client/single server configuration
- ▶ **Three-tiered:** each layer on separate machine

Traditional two-tiered configurations



Being client and server at the same time

Three-tiered architecture



Alternative organizations

Vertical distribution

Comes from dividing distributed applications into three logical layers, and running the components from each layer on a different server (machine).

Horizontal distribution

A client or server may be physically split up into logically equivalent parts, but each part is operating on its own share of the complete data set.

Peer-to-peer architectures

Processes are all equal: the functions that need to be carried out are represented by every process $\Rightarrow$ each process will act as a client and a server at the same time (i.e., acting as a **servant**).

Structured P2P


Essence

Make use of a **semantic-free index**: each data item is uniquely associated with a key, in turn used as an index. Common practice: use a **hash function**

$$\text{key}(\text{data item}) = \text{hash}(\text{data item's value}).$$

P2P system now responsible for storing $(\text{key}, \text{value})$ pairs.

Simple example: hypercube




Looking up d with **key** $k \in \{0, 1, 2, \dots, 2^4 - 1\}$ means **routing** request to node with **identifier** k .

Example: Chord

Principle

- ▶ Nodes are logically organized in a ring. Each node has an m -bit **identifier**.
- ▶ Each data item is hashed to an m -bit **key**.
- ▶ Data item with key k is stored at node with smallest identifier $id \geq k$, called the **successor** of key k .
- ▶ The ring is extended with various **shortcut links** to other nodes.

Example: Chord



Unstructured P2P

Essence

Each node maintains an ad hoc list of neighbors. The resulting overlay resembles a **random graph**: an edge (u, v) exists only with a certain probability $P[(u, v)]$.

Searching

- ▶ **Flooding**: issuing node u passes request for d to all neighbors. Request is ignored when receiving node had seen it before. Otherwise, v searches locally for d (recursively). May be limited by a **Time-To-Live**: a maximum number of hops.
- ▶ **Random walk**: issuing node u passes request for d to randomly chosen neighbor, v . If v does not have d , it forwards request to one of *its* randomly chosen neighbors, and so on.

Flooding versus random walk

Model

Assume N nodes and that each data item is replicated across r randomly chosen nodes.

Random walk

$P[k]$ probability that item is found after k attempts:

$$P[k] = \frac{r}{N} \left(1 - \frac{r}{N}\right)^{k-1}.$$

S ("search size") is expected number of nodes that need to be probed:

$$S = \sum_{k=1}^N k \cdot P[k] = \sum_{k=1}^N k \cdot \frac{r}{N} \left(1 - \frac{r}{N}\right)^{k-1} \approx N/r \text{ for } 1 \ll r \leq N.$$

Flooding versus random walk

Flooding

- ▶ Flood to d randomly chosen neighbors
- ▶ After k steps, some $R(k) = d \cdot (d - 1)^{k-1}$ will have been reached (assuming k is small).
- ▶ With fraction r/N nodes having data, if $\frac{r}{N} \cdot R(k) \geq 1$, we will have found the data item.

Comparison


- ▶ If $r/N = 0.001$, then $S \approx 1000$
- ▶ With flooding and $d = 10, k = 4$, we contact 7290 nodes.
- ▶ Random walks are more communication efficient, but might take longer before they find the result.

Super-peer networks

Essence

It is sometimes sensible to break the symmetry in pure peer-to-peer networks:

- ▶ When searching in unstructured P2P systems, having **index servers** improves performance
- ▶ Deciding where to store data can often be done more efficiently through **brokers**.



Skype's principle operation: A wants to contact B

Both A and B are on the public Internet

- ▶ A TCP connection is set up between A and B for control packets.
- ▶ The actual call takes place using UDP packets between negotiated ports.

A operates behind a firewall, while B is on the public Internet

- ▶ A sets up a TCP connection (for control packets) to a super peer S
- ▶ S sets up a TCP connection (for relaying control packets) to B
- ▶ The actual call takes place through UDP and directly between A and B


Both A and B operate behind a firewall

- ▶ A connects to an online super peer S through TCP
- ▶ S sets up TCP connection to B .
- ▶ For the actual call, another super peer is contacted to act as a **relay** R : A sets up a connection to R , and so will B .

Edge-server architecture

Essence


Systems deployed on the Internet where servers are placed **at the edge** of the network: the boundary between enterprise networks and the actual Internet.



Collaboration: The BitTorrent case

Principle: search for a file F

- ▶ Lookup file at a global directory $\Rightarrow$ returns a **torrent file**
- ▶ Torrent file contains reference to **tracker**: a server keeping an accurate account of **active** nodes that have (chunks of) F .
- ▶ P can join **swarm**, get a chunk for free, and then trade a copy of that chunk for another one with a peer Q also in the swarm.



BitTorrent under the hood

Some essential details

- ▶ A tracker for file F returns the set of its downloading processes: the current **swarm**.
- ▶ A communicates only with a subset of the swarm: the **neighbor set** N_A .
- ▶ if $B \in N_A$ then also $A \in N_B$.
- ▶ Neighbor sets are regularly updated by the tracker

Exchange blocks

- ▶ A file is divided into equally sized **pieces** (typically each being 256 KB)
- ▶ Peers exchange **blocks** of pieces, typically some 16 KB.
- ▶ A can upload a block d of piece D , only if it has piece D .
- ▶ Neighbor B belongs to the **potential set** P_A of A , if B has a block that A needs.
- ▶ If $B \in P_A$ and $A \in P_B$: A and B are in a position that they can **trade** a block.

BitTorrent phases

Bootstrap phase

A has just received its first piece (through **optimistic unchoking**: a node from N_A unselfishly provides the blocks of a piece to get a newly arrived node started).

Trading phase

$|P_A| > 0$: there is (in principle) always a peer with whom A can trade.

Last download phase

$|P_A| = 0$: A is dependent on newly arriving peers in N_A in order to get the last missing pieces. N_A can change only through the tracker.

BitTorrent phases

Development of $|P|$ relative to $|N|$.

